

OS 2022 Class Problem Sheet #1

Problem 1.1: *simple cat (scat) using library and system calls*

(6 points)

Write a program `scat` (simple cat or slow cat) that copies data from the standard input to the standard output.

- a) Implement the data copying loop using the C library functions `getc()/putc()` and using the system calls `read()/write()`, copying on a single byte in each iteration. Your `scat` program should accept the command line options `-l` and `-s`: The option `-l` selects the C library copy loop while the option `-s` selects the system call copy loop. In case there are multiple options on the command line, the last option wins. If there is neither a `-l` nor a `-s` option, the program uses the C library copy loop.
- b) Use your `scat` program to copy a large file to `/dev/null` (a device file that discards all data) and measure the execution times:

```
time ./scat -l < some-large-file > /dev/null
time ./scat -s < some-large-file > /dev/null
```

Repeat the measurements a few times to get stable results. What do you observe? Explain. Use `strace` to investigate the read/write sizes that are used by the two variants of your program. How many read/write calls in total are executed while copying your large file?

- c) Implement another copy loop that uses the Linux specific `sendfile()` system call. The `-p` option selects this copy loop. Set the amount of data that is copied in each call of `sendfile()` such that it matches the amount of bytes read and written by the C library copy loop. Measure the execution time:

```
time ./scat -p < some-large-file > /dev/null
```

What do you observe? Explain.

Hand in the source code of your `scat` program and the results of your analysis. Make sure that your program handles *all* error situations appropriately. Use the `getopt()` function of the C library for parsing command line options.