


Probabilistic Localization

Relative Localization

- dead-reckoning / odometry, etc.
- accumulation of error

⇒ how to

- keep track?
- correct when possible?
(using global localization feedback)




Localization Error Representation

- 3-dim Gaussian for pose
 - expected x , y , θ
 - plus covariances
- noisy motion estimates increase uncertainty

note: approximate model

- ignores non-Gaussian noise,
- e.g., bump noise




example for x, y

Gaussian


Gaussian (aka Normal) distribution

- mean μ (aka expectation)
- variance σ^2 (standard deviation σ)

$$f_{\text{Gaussian}}(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$
$$= N(\mu, \sigma)$$




Gaussian



Multivariate Gaussian

- distribution over k random variables $x = (x_1, \dots, x_k)$
- mean vector μ , covariance matrix Σ

$$f_{\text{Gaussian}}(x) = (2\pi)^{-\frac{k}{2}} |\Sigma|^{-\frac{1}{2}} e^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)} = N(\mu, \Sigma)$$





notes:

- $|A| = \det(A)$ (determinant)
- robotics, often covariance as C (not Σ)
- confidence intervals become confidence regions in form of (hyper)ellipsoids

Localization Error Visualization

(2D) localization error
= Gaussian in (x, y, θ)

- just consider x, y for display
- equidensity contours
 - contours with equal prob. mass
 - are ellipsoids for Gaussian
- principal axes
 - directions = eigenvectors of Σ
 - squared (relative) lengths = corresponding eigenvalues



Error Propagation

more precisely, uncertainty propagation, i.e.,

- development of covariance C
- of random vector x under function $f()$

case 1: linear fct $f()$, i.e., $f(x)=Ax$

$$C^f = ACA^T$$

case 2: arbitrary fct $f()$ with Jacobian J_f

(note: $J_A = A$)

$$C^f = J_f C J_f^T$$

Error Propagation

covariance C , fct $f()$ with Jacobian J_f


$$C^f = J_f C J_f^T$$

- aka *error propagation rule*
or even *error propagation law*
(though more “dirty hack” than “law” 😊)
- based on first order Taylor approximation
- i.e., linearization at point c $f(x) \approx f(c) + J_f(x - c)$

Error Propagation

e.g., 1-dim: consider $[\mu - \sigma, \mu + \sigma]$

$$Y = f(X) \approx f(\mu_x) + \left. \frac{\partial f}{\partial X} \right|_{X=\mu_x} (X - \mu_x)$$




$$\mu_Y = f(\mu_X)$$

$$\sigma_Y = \left. \frac{\partial f}{\partial X} \right|_{X=\mu_x} \sigma_X$$

Example: Differential Drive Odometry

recap: incremental time updates Δt

- $d_{l/r}$: distance by left/right wheel in Δt ($d_{l/r} = v_{l/r} \Delta t$)
- robot drives arc, but Δt small: $\Delta d \approx \Delta s$



motion in Δt


- $\Delta d = \frac{1}{2} (d_l + d_r)$
- $\Delta\theta = (d_r - d_l) / D$

(D = distance of wheels)

Example: Differential Drive Odometry

pose update:

- $x_{t+1} = x_t + \Delta x$
- $y_{t+1} = y_t + \Delta y$
- $\theta_{t+1} = \theta_t + \Delta\theta$
- $\Delta x = \cos(\theta + \Delta\theta/2) \Delta d$
- $\Delta y = \sin(\theta + \Delta\theta/2) \Delta d$
- $\Delta\theta = (d_l - d_r)/D$



Example: Differential Drive Odometry

- $x_{t+1} = x_t + \cos(\theta + \Delta\theta/2) \Delta d$
- $y_{t+1} = y_t + \sin(\theta + \Delta\theta/2) \Delta d$
- $\theta_{t+1} = \theta_t + \Delta\theta$
- $d_{l/r}$: dist. by left/right wheel
- $\Delta d = \frac{1}{2} (d_l + d_r)$
- $\Delta\theta = (d_l - d_r)/D$

$$f(x, y, \theta, d_l, d_r) = (x, y, \theta)^T + (\Delta x, \Delta y, \Delta\theta)^T$$

$$= \begin{pmatrix} x \\ y \\ \theta \end{pmatrix} + \begin{pmatrix} \frac{d_r + d_l}{2} \cos\left(\theta + \frac{d_r - d_l}{D}\right) \\ \frac{d_r + d_l}{2} \sin\left(\theta + \frac{d_r - d_l}{D}\right) \\ \frac{d_r - d_l}{D} \end{pmatrix}$$

Example: Differential Drive Odometry

$$f(x, y, \theta, d_l, d_r) = \begin{pmatrix} x + \frac{d_r + d_l}{2} \cos\left(\theta + \frac{d_r - d_l}{2D}\right) \\ y + \frac{d_r + d_l}{2} \sin\left(\theta + \frac{d_r - d_l}{2D}\right) \\ \theta + \frac{d_r - d_l}{D} \end{pmatrix}$$

uncertainty update:

$$C_{t+1} = \underbrace{J_{f(x,y,\theta)} C_t^{(x,y,\theta)} J_{f(x,y,\theta)}^T}_{\text{motion component}} + \underbrace{J_{f(d_l,d_r)} C_t^{(\Delta x, \Delta y, \Delta \theta)} J_{f(d_l,d_r)}^T}_{\text{wheel-slip}}$$

Example: Differential Drive Odometry

$$f(x, y, \theta, d_l, d_r) = \begin{pmatrix} x + \frac{d_r + d_l}{2} \cos\left(\theta + \frac{d_r - d_l}{2D}\right) \\ y + \frac{d_r + d_l}{2} \sin\left(\theta + \frac{d_r - d_l}{2D}\right) \\ \theta + \frac{d_r - d_l}{D} \end{pmatrix} \quad J_{f(x, y, \theta)} = \begin{pmatrix} \frac{\partial f}{\partial x} & \frac{\partial f}{\partial y} & \frac{\partial f}{\partial \theta} \end{pmatrix}$$

$$= \begin{pmatrix} 1 & 0 & -\frac{d_r + d_l}{2} \sin\left(\theta + \frac{d_r - d_l}{2D}\right) \\ 0 & 1 & \frac{d_r + d_l}{2} \cos\left(\theta + \frac{d_r - d_l}{2D}\right) \\ 0 & 0 & 1 \end{pmatrix}$$

$$= \begin{pmatrix} 1 & 0 & -\Delta d \sin(\theta + \Delta\theta/2) \\ 0 & 1 & \Delta d \cos(\theta + \Delta\theta/2) \\ 0 & 0 & 1 \end{pmatrix}$$

Example: Differential Drive Odometry

$$f(x, y, \theta, d_l, d_r) = \begin{pmatrix} x + \frac{d_r + d_l}{2} \cos\left(\theta + \frac{d_r - d_l}{2D}\right) \\ y + \frac{d_r + d_l}{2} \sin\left(\theta + \frac{d_r - d_l}{2D}\right) \\ \theta + \frac{d_r - d_l}{D} \end{pmatrix}$$

$$\begin{aligned} J_{f(d_l, d_r)} &= \begin{pmatrix} \frac{\partial f}{\partial d_l} & \frac{\partial f}{\partial d_r} \end{pmatrix} \\ &= \begin{pmatrix} \frac{1}{2} \cos\left(\theta + \frac{d_r - d_l}{2D}\right) + \frac{d_r + d_l}{4D} \sin\left(\theta + \frac{d_r - d_l}{2D}\right) & \frac{1}{2} \cos\left(\theta + \frac{d_r - d_l}{2D}\right) - \frac{d_r + d_l}{4D} \sin\left(\theta + \frac{d_r - d_l}{2D}\right) \\ -\frac{1}{2} \sin\left(\theta + \frac{d_r - d_l}{2D}\right) - \frac{d_r + d_l}{4D} \cos\left(\theta + \frac{d_r - d_l}{2D}\right) & \frac{1}{2} \sin\left(\theta + \frac{d_r - d_l}{2D}\right) + \frac{d_r + d_l}{4D} \cos\left(\theta + \frac{d_r - d_l}{2D}\right) \\ -\frac{1}{D} & \frac{1}{D} \end{pmatrix} \\ &= \begin{pmatrix} \frac{1}{2} \cos(\theta + \Delta\theta/2) + \frac{\Delta d}{2D} \sin(\theta + \Delta\theta/2) & \frac{1}{2} \cos(\theta + \Delta\theta/2) - \frac{\Delta d}{2D} \sin(\theta + \Delta\theta/2) \\ -\frac{1}{2} \sin(\theta + \Delta\theta/2) - \frac{\Delta d}{2D} \cos(\theta + \Delta\theta/2) & \frac{1}{2} \sin(\theta + \Delta\theta/2) + \frac{\Delta d}{2D} \cos(\theta + \Delta\theta/2) \\ -\frac{1}{D} & \frac{1}{D} \end{pmatrix} \end{aligned}$$

Example: Differential Drive Odometry

$$\mathbf{C}_{t+1} = \mathbf{J}_{f(x,y,\theta)} \mathbf{C}_t^{(x,y,\theta)} \mathbf{J}_{f(x,y,\theta)}^T + \mathbf{J}_{f(d_l,d_r)} \mathbf{C}_t^{(\Delta x,\Delta y,\Delta \theta)} \mathbf{J}_{f(d_l,d_r)}^T$$

- start
 - just start pose at $(0,0,0)^T$
 - and perfect localization
- modeling of slip
 - proportional to left/right wheel distance
 - with constant(s), e.g., found via calibration

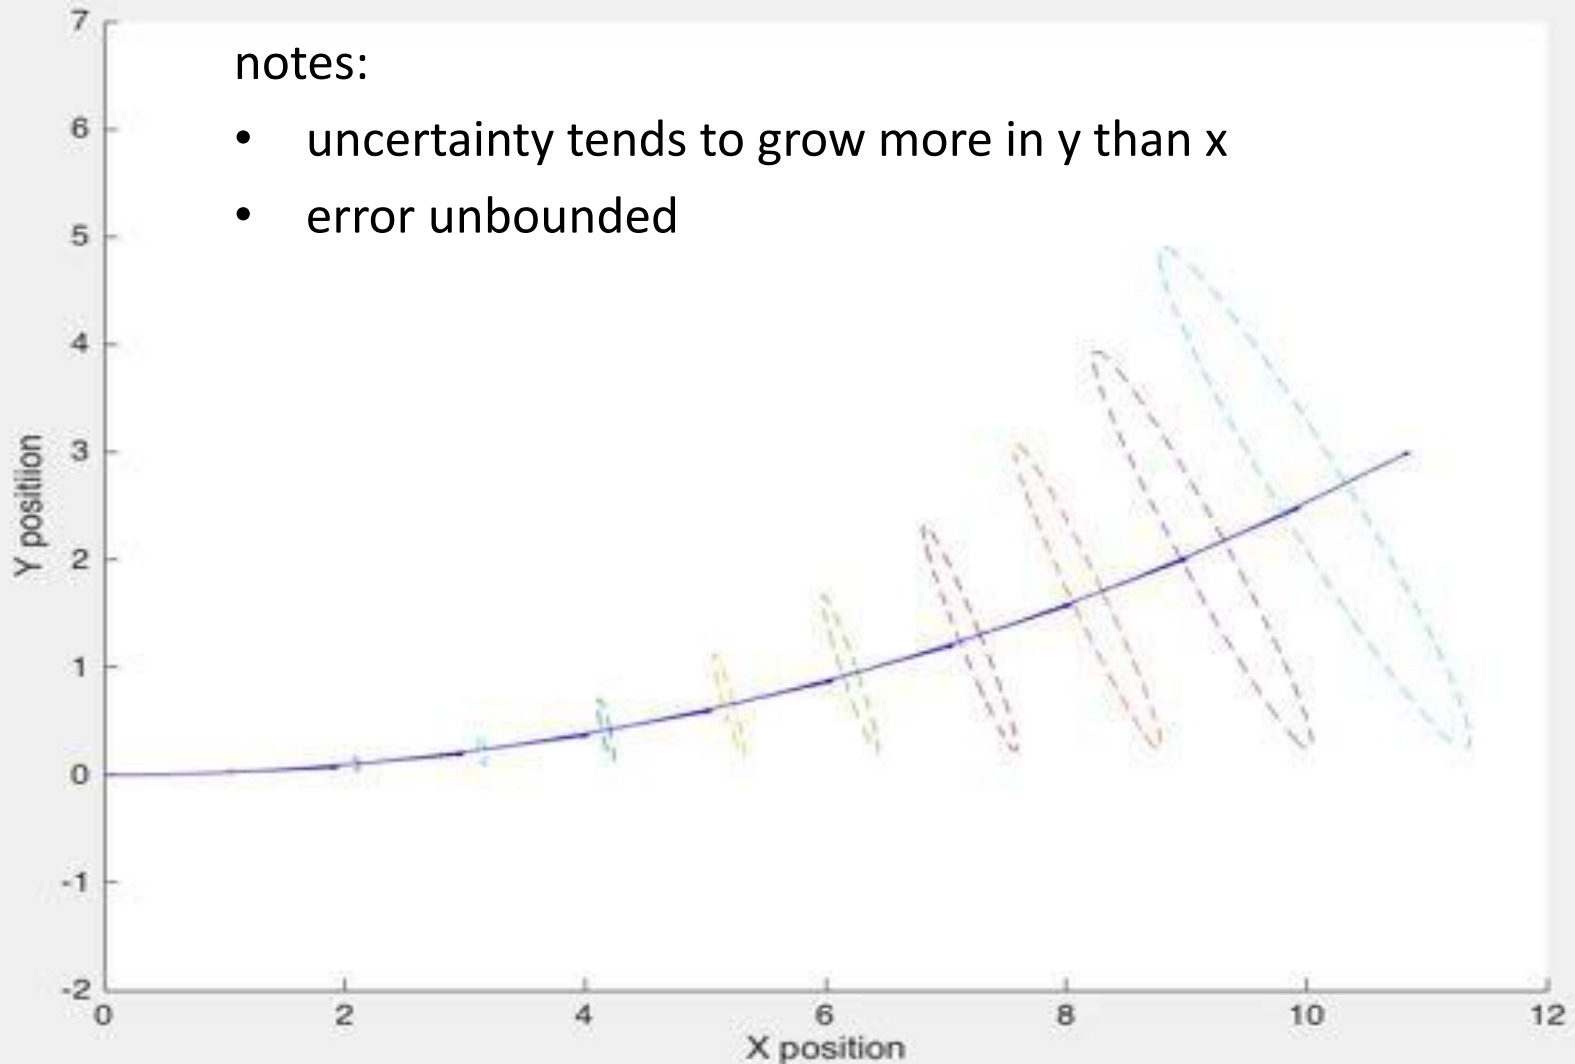
$$\mathbf{C}_0^{(x,y,\theta)} = \mathbf{0}$$

$$\mathbf{C}_t^{(\Delta x,\Delta y,\Delta \theta)} = \begin{pmatrix} c_l |d_l| & 0 \\ 0 & c_r |d_r| \end{pmatrix}$$

Example: Differential Drive Odometry

notes:

- uncertainty tends to grow more in y than x
- error unbounded



Kalman Filter


Kalman Filter

- **recursive** data processing algorithm
 - no need to store all previous measurements
 - and to reprocess all data at each time step
- to generate **optimal** estimate from measurements
 - for linear system (but also widely used for non-linear, too)
 - and white Gaussian noise (i.e., uncorrelated in time)

can e.g., be used for localization

- e.g., by incorporating (noisy) absolute measurements
- in (noisy) relative motion estimates (by e.g., odometry)


Simple Example: 1D Motion



- 1-dimensional localization: position $x(t)$
- odometry & perception of a landmark
- assume Gaussian distributed measurements

Simple Example: 1D Motion


1st: **standing still** and multiple, noisy measurements



- location estimate at t_1 : mean $\mu_1 = z_1$ and variance $= \sigma^2_1$
- first estimate of position: $\hat{x}(t_1) = \mu_1 = z_1$
- error variance of estimate: $\sigma^2(t_1) = \sigma^2_1$
- robot in same position at time t_2 : predicted position is z_1

Simple Example: 1D Motion

1st: **standing still** and multiple, noisy measurements



- prediction $\hat{y}^-(t_2)$
- landmark measurement at t_2 : mean = z_2 and variance = σ^2_2
- correct prediction with this measurement to get $\hat{x}(t_2)$
- via linear interpolation with variances as weights

Fusing the data

- interpolation is "best" combination
 - based on statistical criteria, namely
 - maximum likelihood estimate and
 - minimum variance of all possible linear combinations

$$\mu = \frac{\sigma_2^2}{\sigma_1^2 + \sigma_2^2} z_1 + \frac{\sigma_1^2}{\sigma_1^2 + \sigma_2^2} z_2$$


$$\sigma^2 = \left(\frac{1}{\sigma_1^2} + \frac{1}{\sigma_2^2} \right)^{-1}$$

- new state estimate then $\hat{x}(t_2) = \mu$

The influence of the variances

- variances σ_i^2 determine
 - how much to trust the measurements
- e.g., variances are equal
 - => sensor measurements are averaged
- one is larger, e.g., σ_1^2
 - there is more uncertainty in the measurement z_1
 - and it is weighted less heavily than z_2
- so, even poor quality data contains information
 - it is included without "spoiling" the better quality data
 - and it improves the output of the filter

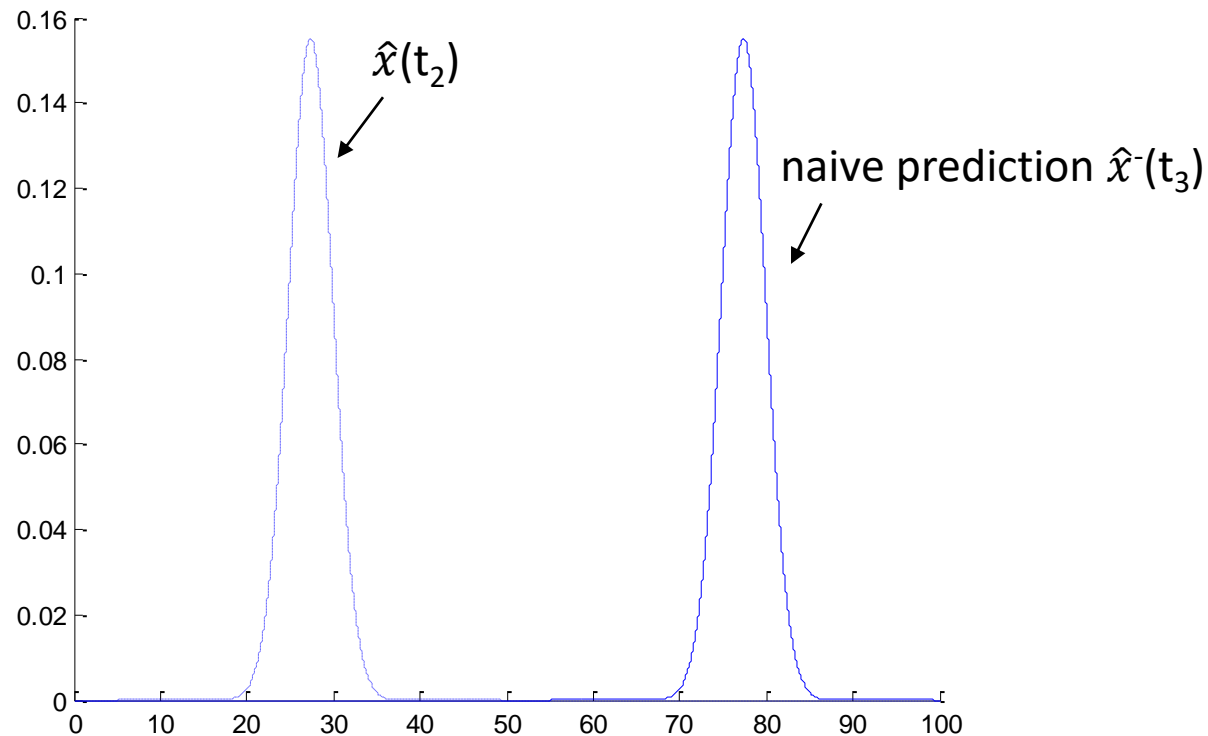
Simple Example: 1D Motion



- corrected mean: new optimal estimate of position
- new variance is smaller than either of the previous two variances


Simple Example: 1D Motion

2nd: **including noisy motion** (and related uncertainty based model)




- at time t_3 , robot moves with velocity $dx/dt=u$
- naive approach: shift probability to the right, i.e., add ut to mean
- but motion estimates (odometry) are noisy

Simple Example: 1D Motion



- use proper model by adding Gaussian noise w
- $dx/dt = u + w$
- distribution for prediction “moves” and “spreads out”

Simple Example: 1D Motion



- new measurement of the landmark at t_3
- correct again the prediction
- using difference between prediction and measurement aka ***residual*** and so on...

Overview Kalman

- initial conditions (\hat{x}_{k-1} and σ_{k-1})
- prediction (\hat{x}_k^-, σ_k^-)
 - use of initial conditions and model (e.g., odometry)
 - to make prediction
- measurement (z_k)
 - take measurement
- correction (\hat{x}_k, σ_k)
 - use measurement to correct prediction
 - by ‘fusing’ prediction and residual
 - i.e, by merging two Gaussians
 - result: optimal estimate with smaller variance

Kalman Filter

given: linear system with white Gaussian noise

$$x_k = Ax_{k-1} + Bu_{k-1} + w_{k-1}$$

$$z_k = Hx_k + v_k$$

Zero-mean Gaussians with
covariance matrices Q , R

$$p(w) = N(0, Q)$$

$$p(v) = N(0, R)$$

Kalman Filter

- a priori state estimate \hat{x}_k^-
 - at step k
 - includes all knowledge of the process prior to step k
- a posteriori state estimate \hat{x}_k
 - at step k
 - given the current measurement
- a priori and a posteriori estimate errors

$$e_k^- = x_k - \hat{x}_k^- \qquad e_k = x_k - \hat{x}_k$$

- a priori and a posteriori estimate error covariances

$$P_k^- = E(e_k^- e_k^{-T}) \qquad P_k = E(e_k e_k^T)$$

Kalman Filter

blending factor

aka (Kalman) gain



$$\hat{x}_k = \hat{x}_k^- + K_k (z_k - H\hat{x}_k^-)$$

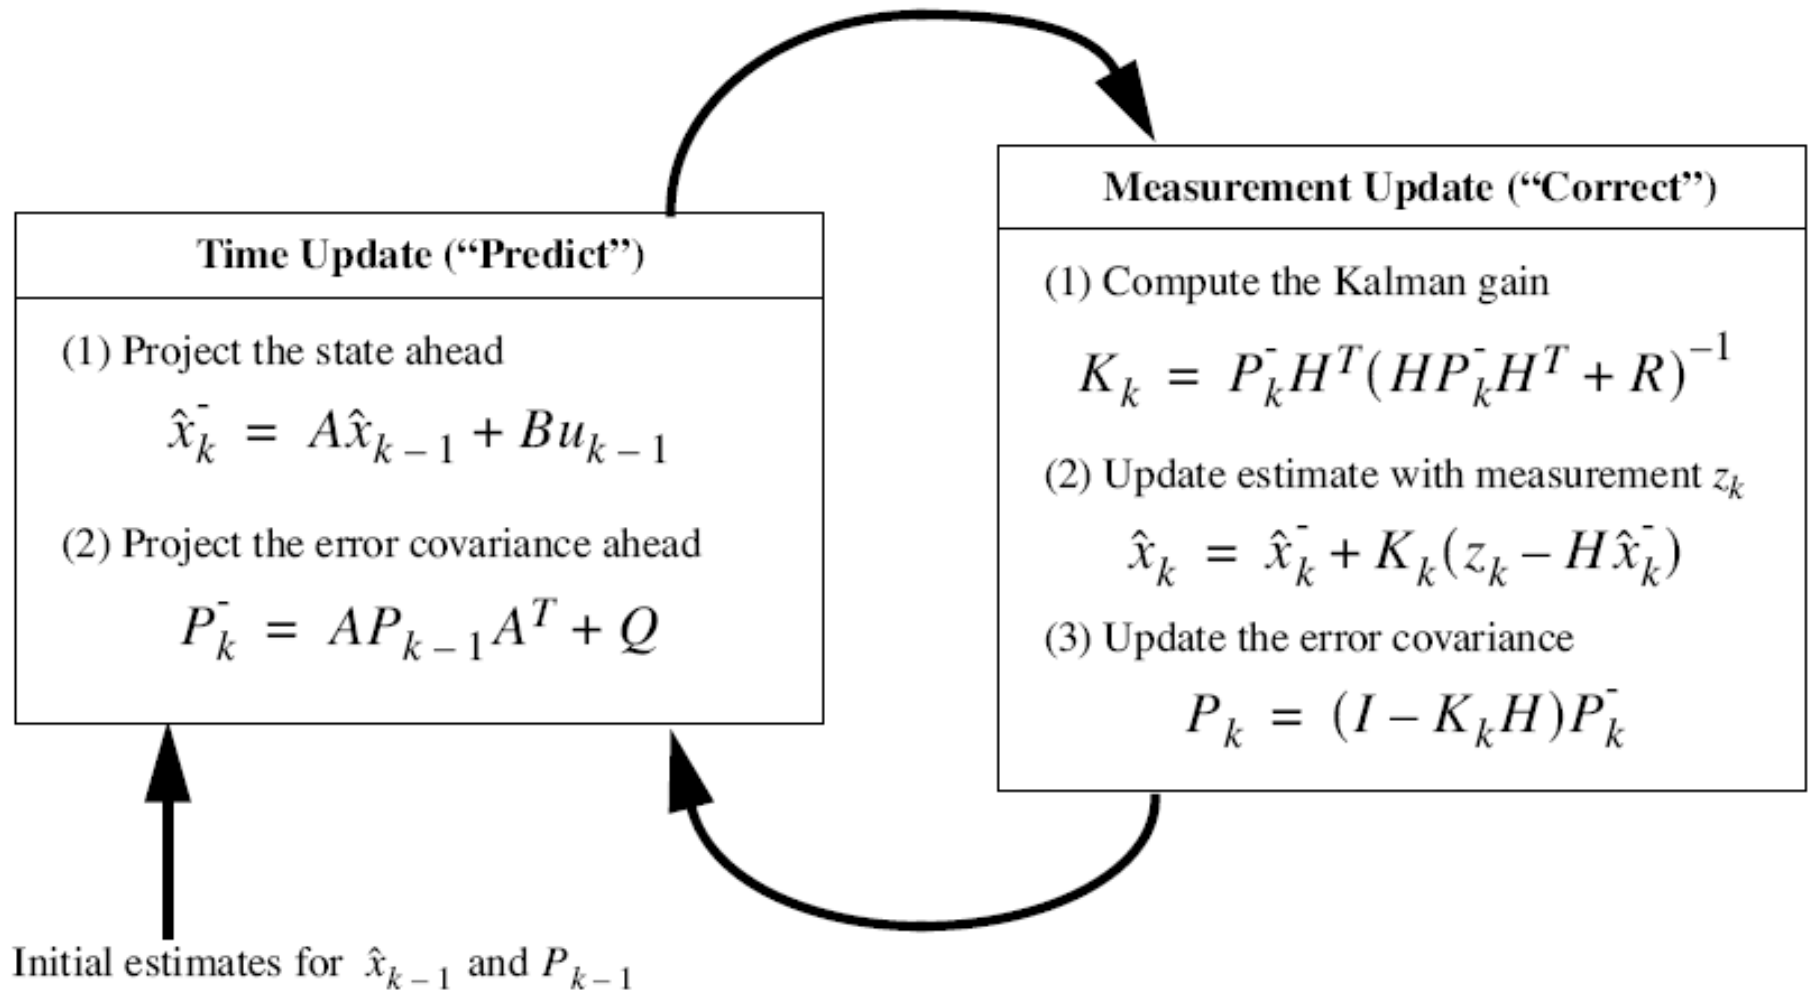


(measurement) innovation

aka residual

$$K_k = P_k^- H^T (H P_k^- H^T + R)^{-1}$$

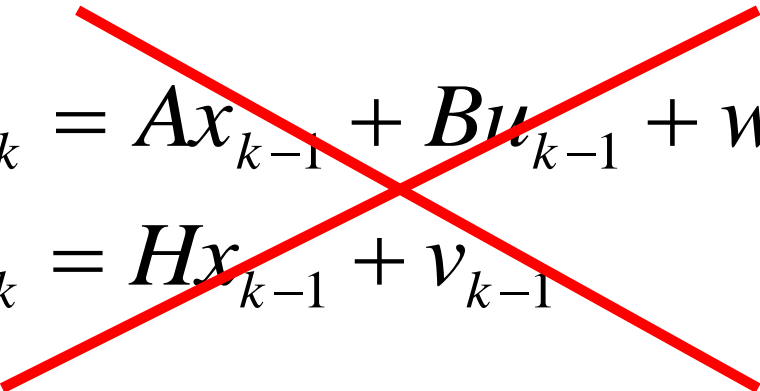
Kalman Filter



Extended Kalman Filter

Extended Kalman Filter

state-evolution and measurement equations
are often non-linear


$$\begin{aligned}x_k &= Ax_{k-1} + Bu_{k-1} + w_{k-1} \\z_k &= Hx_{k-1} + v_{k-1}\end{aligned}$$

⇒ Extended Kalman Filter (EKF)

Extended Kalman Filter

state-evolution and measurement
fct's f and h are non-linear:

$$x_k = f(x_{k-1}, u_{k-1}) + w_{k-1}$$

$$z_k = h(x_k) + v_k$$

Zero-mean Gaussians with
covariance matrices Q , R

$$p(w) = N(0, Q)$$

$$p(v) = N(0, R)$$

Extended Kalman Filter

linearization of update equations Jacobian of

- $f : J_f$
- $h : J_h$

- predictor step: $\hat{x}_k^- = f(\hat{x}_{k-1}, u_{k-1})$

$$P_k^- = J_f P_{k-1} J_f^T + Q$$

- Kalman gain: $K_k = P_k^- J_h^T (J_h P_k^- J_h^T + R)^{-1}$

$$\hat{x}_k = \hat{x}_k^- + K_k (z_k - h(\hat{x}_k^-))$$

- corrector step:

$$P_k = (I - K_k J_h) P_k^-$$

Extended Kalman Filter

- Extended Kalman Filter straightforward to use
- but unfortunately known to be often not stable,
- i.e., to diverge due to the linearization


=> Unscented Kalman Filter (UKF)

Unscented Kalman Filter

Unscented Kalman Filter (UKF)

basic idea:

- do not linearize transformation
- but choose (few) sample points
- to represent mean and covariance




Unscented Kalman Filter (UKF)

basic idea: choose (few) sample points for mean and covariance

advantages


- (can be) more accurate than EKF
- no need for Jacobians




Unscented Kalman Filter (UKF)



set of sigma points




transform each one
with non-linear fct




compute new Gaussian from
transformed and weighted points

Comparison EKF and UKF


Actual (sampling)



Linearized (EKF)



UT



Unscented Kalman Filter

- sigma points χ_i , weights w_i
- choose so that

$$\sum_i w_i = 1$$

$$\mu = \sum_i w_i \chi_i$$

$$\Sigma = \sum_i w_i (\chi_i - \mu)(\chi_i - \mu)^T$$

note: there is no unique solution for χ_i , w_i
(there are several versions of UKF)

Unscented Kalman Filter

n : dimension

λ : scaling parameter

$$\chi_0 = \mu$$

$$\chi_i = \mu + \left(\sqrt{(n + \lambda)\Sigma} \right)_i \quad \text{for } i = 1, \dots, n$$

$$\chi_i = \mu - \left(\sqrt{(n + \lambda)\Sigma} \right)_{i-n} \quad \text{for } i = n + 1, \dots, 2n$$

note: $(A)_i$ = column vector i of A

Matrix Square Root

matrix square root: S with $SS^T = \Sigma$

recap:

- Σ symmetric (hence, $c\Sigma$ symmetric)
- SVD of symmetric matrix $\Sigma = VLV^T$

$$\Sigma = SS^T = VLV^T$$

$$\Rightarrow S = VL^{(1/2)}$$

Matrix Square Root

matrix square root: S with $SS^T = \Sigma$

Cholesky decomposition

- $\Sigma = LDL^T = LD^{1/2}(D^{1/2}L)^T = GG^T$
 - symmetric Σ
 - D : positive diagonal matrix
 - L : normed lower triangular matrix
(normed: diagonal is all 1's)
 - G : lower triangular matrix
- often used in UKF implementations

Unscented Kalman Filter

computing the weights

- w^m : for mean
- w^c : for covariance

$$w_0^m = \frac{\lambda}{n + \lambda}$$

$$w_0^c = w_0^m + (1 - \alpha^2 + \beta)$$

$$w_i^m = w_i^c = \frac{1}{2(n + \lambda)} \quad \text{for } i = 1, \dots, 2n$$

α, β : parameters

Unscented Kalman Filter

$$\chi_0 = \mu$$

$$\chi_i = \mu + \left(\sqrt{(n + \lambda) \Sigma} \right)_i \quad (i \leq n)$$

$$\chi_i = \mu - \left(\sqrt{(n + \lambda) \Sigma} \right)_{i-n} \quad (i > n)$$

$$w_0^m = \frac{\lambda}{n + \lambda}$$

$$w_0^c = w_0^m + (1 - \alpha^2 + \beta)$$

$$w_i^m = w_i^c = \frac{1}{2(n + \lambda)}$$

parameters: $\alpha \in]0,1]$

$\beta = 2$ (optimal for Gaussians)

$$\lambda = \alpha^2 (n + \kappa) - n$$

$$\kappa \geq 0$$

Unscented Kalman Filter

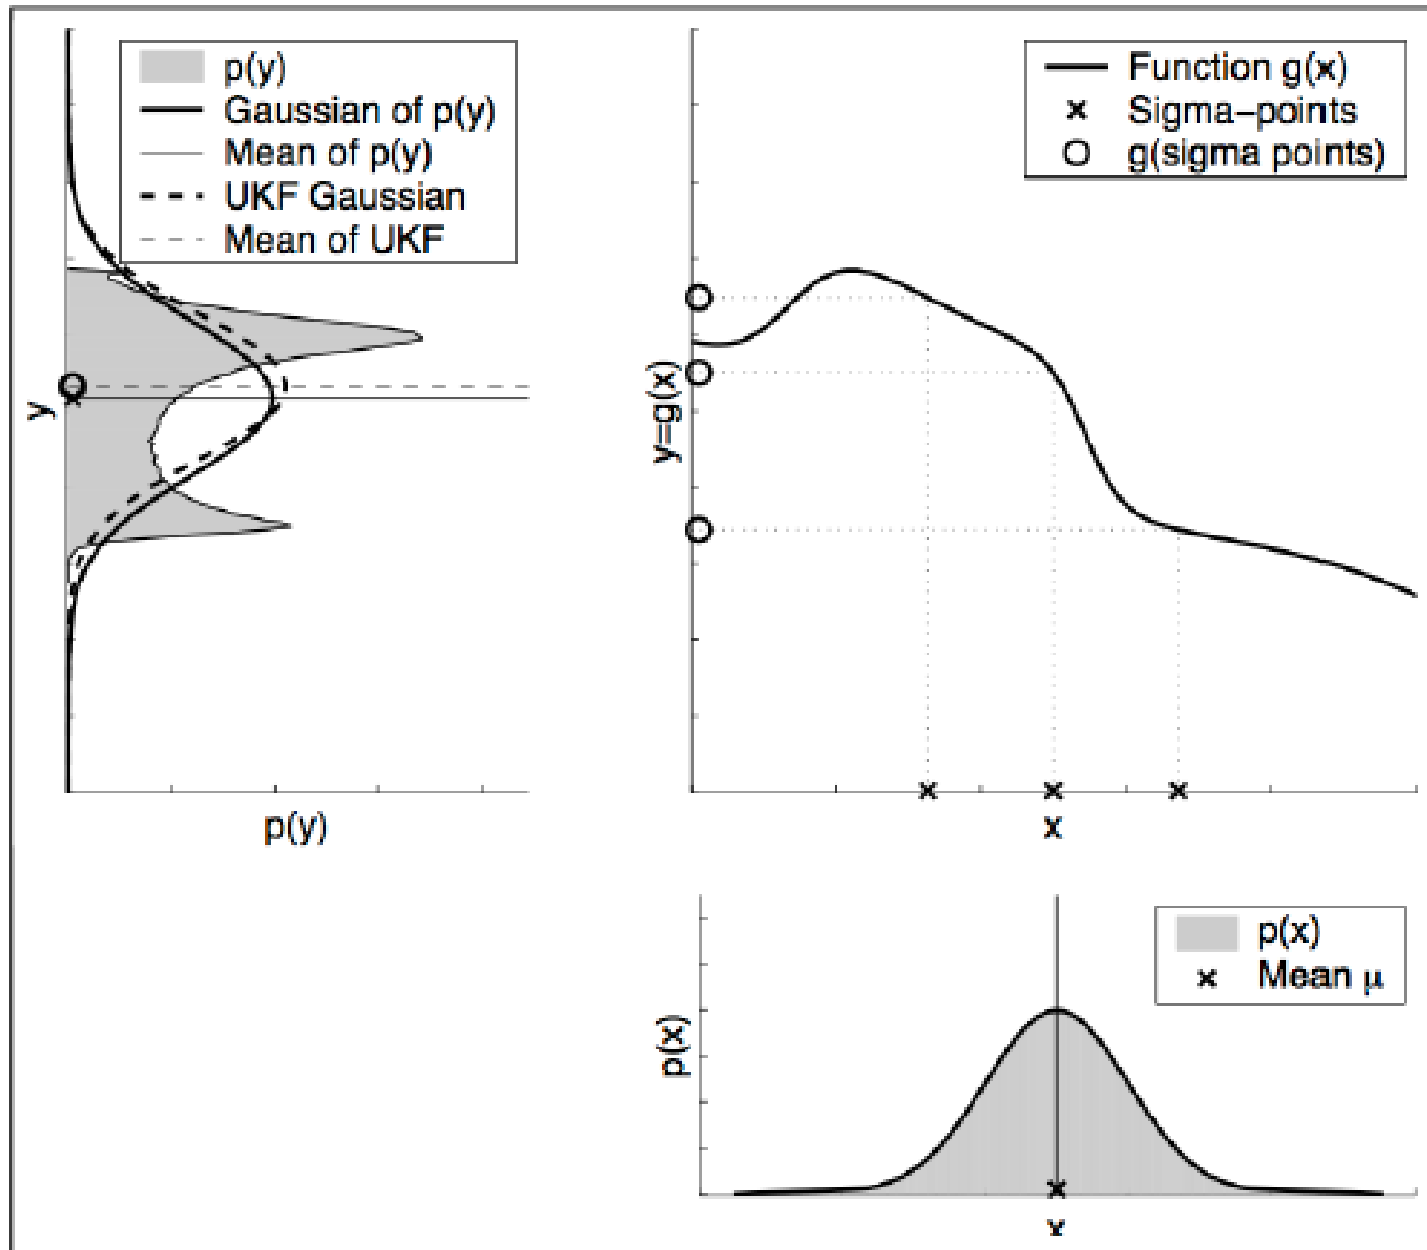
new Gaussian $N(\mu', \Sigma')$

from sigma points transformed by non-linear $f()$

$$\mu' = \sum_{i=0}^{2n} w_i^m f(\chi_i)$$

$$\Sigma' = \sum_{i=0}^{2n} w_i^c (f(\chi_i) - \mu')(f(\chi_i) - \mu')^T$$

Example



Unscented Kalman Filter

predict:

$$\chi_{i,k} = f(\chi_{i,k-1})$$

$$z_{i,k-1} = h(\chi_{i,k-1})$$

$$\hat{x}_k^- = \sum_{i=0}^{2n} w_i^m \chi_{i,k-1}$$

$$\hat{z}_k^- = \sum_{i=0}^{2n} w_i^m z_{i,k-1}$$

$$P_k^- = \sum_{i=0}^{2n} w_i^c (\chi_{i,k} - \hat{x}_k^-)(\chi_{i,k} - \hat{x}_k^-)^T$$

correct:

$$P_{z_k z_k} = \sum_{i=0}^{2n} w_i^c (z_{i,k} - \hat{z}_k^-)(z_{i,k} - \hat{z}_k^-)^T$$

$$P_{x_k z_k} = \sum_{i=0}^{2n} w_i^c (\chi_{i,k} - \hat{x}_k^-)(z_{i,k} - \hat{z}_k^-)^T$$

$$K_k = P_{x_k z_k} P_{z_k z_k}^{-1}$$

$$\hat{x}_k = \hat{x}_k^- + K_k (z_k - \hat{z}_k^-)$$

$$P_k = P_k^- - K_k P_{z_k z_k} K_k^T$$

Particle Filter

Particle Filter (PF)

- represent distribution
 - by **randomly chosen weighted samples** (particles)
 - population based somewhat similar to Evolutionary Algorithm
- particles are transformed under systems dynamics (model)
- test predicted states (particles) with observation
- do a selection
 - multiply or discard particles
 - i.e., survival of the fittest
- hence „Monte Carlo“ filter (aka „condensing“)
- (likely) convergence depends on #samples, i.e., particles

Particle Filter

example: localization

(from Dieter Fox, UWash)

- red dot
 - particle
 - estimated robot pose
 - init: random
- 24 sonar sensors
 - match range with given map
 - basis for selection

more about PF soon...

