

A Look into Mapping

“Simple” Mapping

given:

A) localization plus

B) environment sensing (e.g., free/occupied)

then:

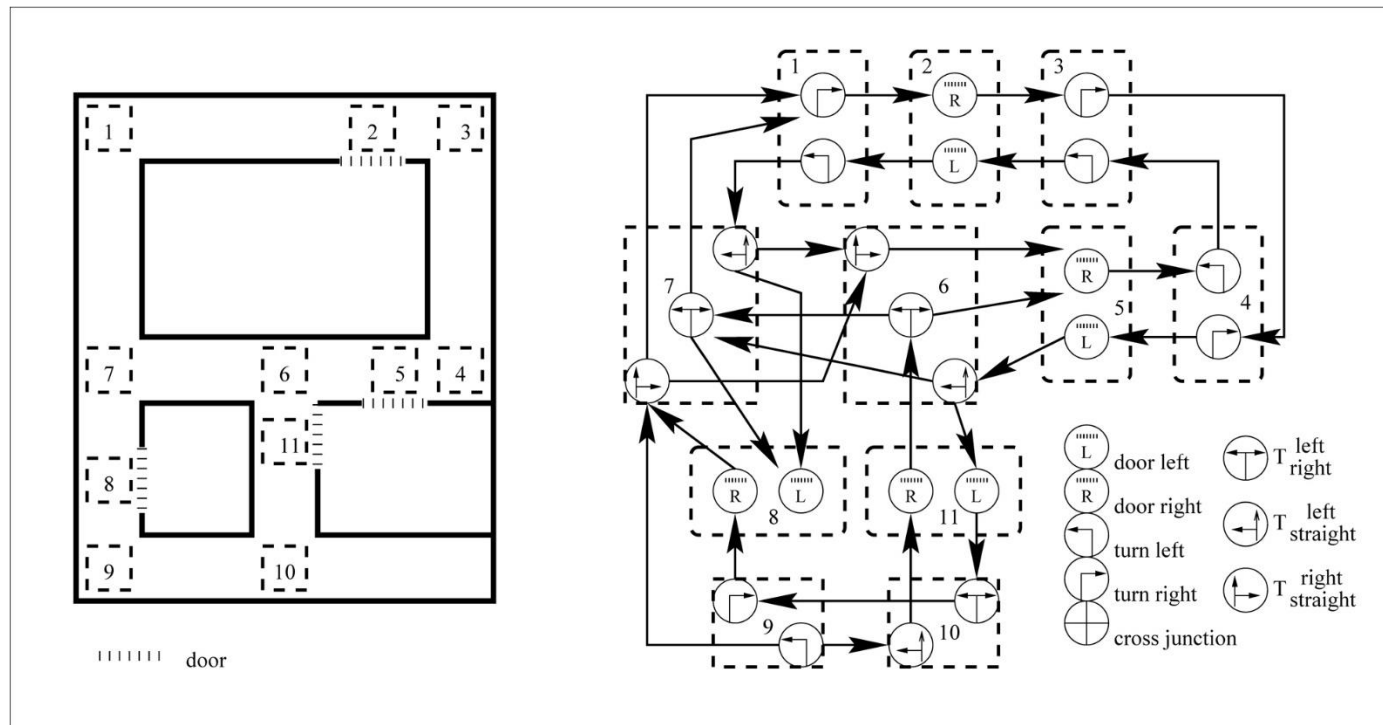
C) enter data into map representation

we now have a (first) look into C), then B)


Short Overview of Map Representations

Maps in general

- geometric (Euclidean)
 - “proper” distances encoded
- topological
 - graph of “places” and “connections”



Metric: 2D/3D Occupancy Grid




- 2D/3D regular grid (array)
 - with occupancy / free space / unknown info in each cell
 - $x, y, (z)$ correspond to metric Cartesian coordinates
- simple but memory intensive
 - predominant representation for 2D
 - 3D less so, alternative: octtree

Metric: 2D Visual (Photo) Data

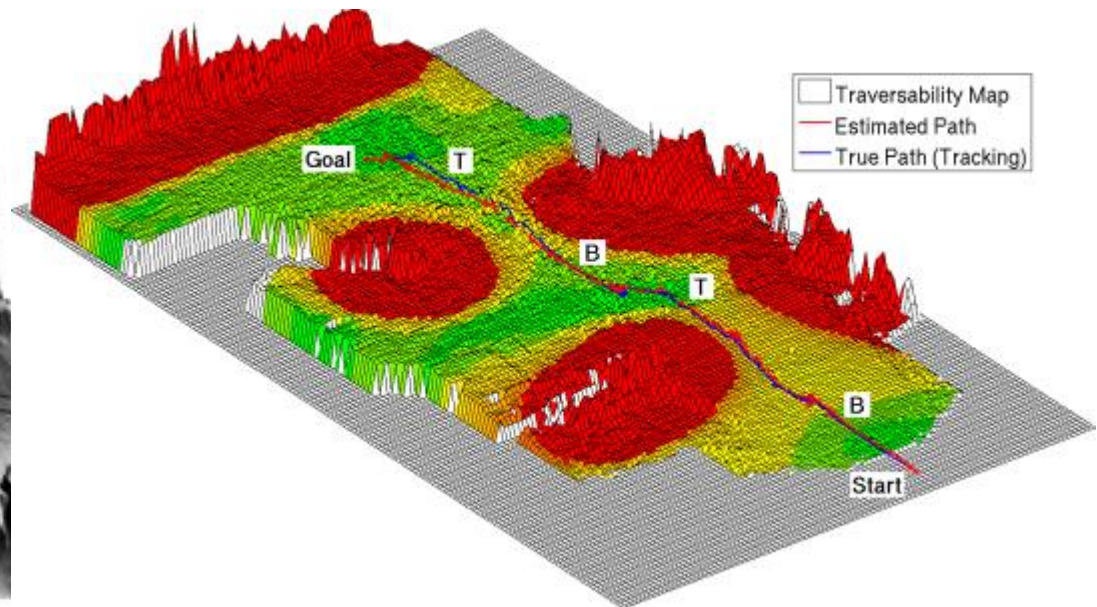

mosaic

- 2D regular grid (array) with visual data in each cell aka **raster image**
- x,y correspond to metric Cartesian coordinates
- geo-referenced: use of geographical coordinates
 - gathered e.g. via GNSS
 - typically just for the (pose of) the origin
 - stored in image (meta-)data; see e.g. GeoPDF, GeoTIFF
- often (with meta-data) in Geo Information System (GIS)



(2.5 D) Elevation Map


- (typically regular) grid: value = elevation
- Digital Elevation Model (DEM),
Digital Terrain Model (DTM),
Digital Surface Model (DSM)
- often falsely denoted as 3D, but 2D manifolds in 3D space



(2.5 D) Elevation Map

bathymetry

- underwater elevation map
- is **not** 3D



3D Point Cloud



- set of 3D points
 - fine grain coordinates
 - point represents occupancy
- pro/cons
 - "raw" range data format
 - can grow very large
 - computational geometrie very hard



2D/3D Quad-/Octree





- volumetric representation
- recursive decomposition
 - divide in 4 / 8 cells
 - each occupied cell is further divided

option for point clouds:
store 3D info in cells



3D Surface Models

- in general
 - volume (voxel) vs surface representation
- surface
 - 3D meshes
 - regular polygons
 - points span surfaces
 - geometric object collections
 - planes, spheres, boxes
 - defined by parameters
- pros/cons
 - compact; relatively few data points needed
 - basis for computer graphics
 - basis for computational geometry
 - supported by standards (e.g., Collada)
 - difficult to extract from raw 3D sensor data



```
Transform {  
  translation 0 0.15 0  
  children [  
    shape {  
      appearance Appearance {  
        material Material {  
          emissivecolor 0 1 0  
        }  
      }  
      geometry cylinder {  
        radius .2  
        height .4  
      }  
    ]  
  }  
}
```

Representing Probabilistic Data: Evidence Grid Map

Conditional Probability

- probability of A under condition B

$$p(A|B) = \frac{p(A \cap B)}{p(B)}$$

- Bayes' theorem

$$p(A|B) = p(B|A) \frac{p(A)}{p(B)}$$

as $p(A \cap B) = p(B \cap A)$

Evidence Grid

grid map alg. (Moravec, Elfes, 80s)

- based on Bayes' theorem

$$p(A | B) = p(B | A) \frac{p(A)}{p(B)}$$

- using odds representation, $p(A)/p(\neg A)$
 - o = cell occupied, s = sensor value
 - grid cell $g(x,y)$: odds of being occupied

$$g(x, y) = \frac{p(o)}{p(\neg o)}$$

Evidence Grid

- we want:
occupancy probability given sensor value $p(o | s)$
- core idea: use sensor model
 - likelihood of getting a sensor value depending on occupancy $p(s | o)$
 - can be derived a priori (through experiments)
- in combination with Bayes' theorem

same principle for free space ($\neg o$)

Incremental Update

- update: Bayes with sensor model

$$\frac{p(o | s)}{p(\neg o | s)} = \frac{p(s | o) \frac{p(o)}{p(s)}}{p(s | \neg o) \frac{p(\neg o)}{p(s)}} = \frac{p(s | o) p(o)}{p(s | \neg o) p(\neg o)} = \frac{p(s | o)}{p(s | \neg o)} \cdot \frac{p(o)}{p(\neg o)}$$

- assume independence of measurements

$$p((o | s_t) \wedge (o | s_{t-1})) = p(o | s_t) \cdot p(o | s_{t-1})$$

$$p((\neg o | s_t) \wedge (\neg o | s_{t-1})) = p(\neg o | s_t) \cdot p(\neg o | s_{t-1})$$

$$g(x, y)_t = \frac{p(s | o)}{p(s | \neg o)} g(x, y)_{t-1}$$

Log Odds for Efficiency


$$g'(x, y) = \log\left(\frac{p(o)}{p(\neg o)}\right) = \log(p(o)) - \log(p(\neg o))$$

- multiplication becomes addition
- simple integer representation (e.g., byte)
- initialization $p(o) = p(\neg o) = 0.5 \Leftrightarrow g'(x, y) = 0$
- sign indicates occupancy, value confidence

Need for sensor models

narrow beam range

- e.g., laser range finder
- here: two aIR sensors
- Bresenham line drawing
 - open space (line): add v1
 - end point: add v2
 - $v1, v2 = \log(p(s|o)/p(s|\neg o))$
 - naive: $v1, v2 = -1, +1$





Need for sensor models

wide beam range

- e.g., ultrasound


- algorithmic approach
 - inconvenient
 - complex distributions

- store profiles
 - lookup tables



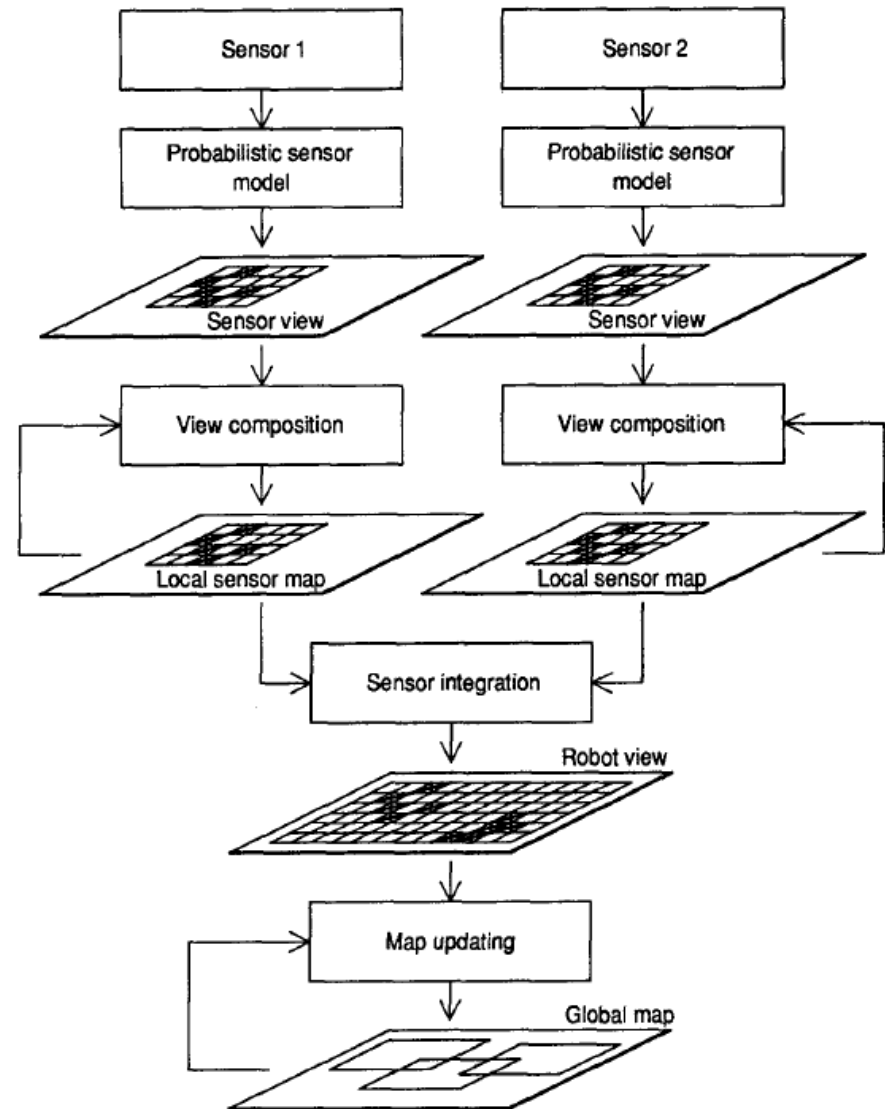
Sensor Models

- unlimited possibilities for modeling
- in theory
 - for all sensor values s
 - distribution of probabilities
 - for all positions x, y (from sensor to global frame) needed
- often in reality
 - simplified coverage and probabilities
 - algorithmic model (e.g. narrow beam) or look up (e.g. ultrasound)



Sensor Fusion

- sensor models
 - data in local grids
- can easily be fused
 - log odds => addition



Range Sensors

(more precisely: distance to obstacles)

Laser Range Finder (LRF)

- aka laser scanner, laser/light radar (ladar/lidar)
- time of flight: speed of light is very fast
 - timing with digital electronics challenging
 - e.g. capacitive discharge for time-measurements
 - interference (slow amplitude modulation)
 - pulsed versions have high sampling-rate
 - rotating mirror: high spatial resolution
- very accurate

LRF Example

- Hokuyo URG-04LX
 - field of view: 240°
 - angular resolution: 0.36°
 - response time: 100 ms
 - resolution: 20 mm
 - range: 4 m
 - power: 2.5 W at 5V
 - weight: 0.16 kg
 - dimensions: 50 x 50 x 70 mm (L x W x H)




LRF Example

- SICK LMS 200
 - field of view: 180°
 - angular resolution: $1 - 0,25^\circ$
 - response time: 13 - 53 ms
 - resolution: 10 mm
 - error: +/- 15 mm
 - range: 80 m
 - power: 20 W at 24V
 - weight: 4,5 kg
 - dimensions: 156 x 155 x 210 mm (L x W x H)



Actuated LRF

- standard LRF: 2D
 - horizontal scan plane
- actuated => 3D
 - servo motor
 - 1 DOF rotation
 - possible movements
 - pitch
 - roll
 - yaw side
 - yaw up



High-End 3D LRF examples


- FARO Focus 3D
 - ~1 million points / sec (color CCD included)
 - 120 m range
 - FOV: 360 x 305 deg
- RIEGL VZ-400
 - 120,000 points / sec
 - 600 m range
 - FOV: 360 x 100 deg
- Velodyne HDL-32E
 - 0.7 million points / sec
 - 32 laser elements
 - FOV: 360 x 40 deg (+10 to -30 deg)
 - 100 m range (+/- 2cm)
 - compact, light, low-power



Structured Light

pattern of light

- with know properties
- typical: grid or pseudo-random
- derive shape
- from the geometric deformations



Structured Light: Kinect

- inexpensive game controller (gestures)
- for MS Xbox 360
 - active IR pattern and sensor
 - plus color camera
 - aka RGBD (color + depth)
- caveats
 - does not work outdoors (IR)
 - field of view and range OK, but...
 - noise and resolution OK, but...



Ultrasound (US) Sensors

- relatively low cost sensors
- field of view
 - often broad (60-120°)
 - narrow ones exist (10-20°)
- limitations
 - speed of sound in air is slow: $\sim 300 \text{ m / s}$
 - and depends on humidity, temperature, ...
 - need to handle echos, multiple paths

US Example

Baumer UNAM 30U6103

- range: 10-70 cm
- resolution: 0.3 mm
- field of view: 10 deg
- power: <1W at 15-30V
- linear analog output 0-10V



typical obstacle sensor,
respectively simple „object“ sensor in automation



US Example

Polaroid 6500 & 600

- ranging board 6500
- transducer 600
- range: up to 10m
- field of view: 120 deg
- power: 5V
 - 2 A during transmit
 - 0.1 A else





Passive Range Sensing

(i.e., vision based)




(Glimpse into) Stereo Vision

- rectified images
 - i.e., no camera distortions and on one plane (to ease search for correspondences: horizontal epipolar line)
- epipolar geometry
 - corresponding projected points (aka epipoles)
 - and center of projection
 - are on a line (aka epipolar line)



Stereo Vision

- epipolar line
 - corresponding points p, p' on a horizontal line
 - distance $|p, p'|$ aka disparity plus triangulation
 - gives range, respectively full 3D coordinates



Stereo Vision

- note: stereo typically used as 2.5D sensor
 - i.e., **dense** feature/range measurements
 - no need for interest points: brute force try all
 - computationally feasible due to epipolar constraint
 - and simple matching (e.g., pixel blocks)
- occasionally (combined with): **sparse** stereo
 - get (few) but very robust 3D landmarks
 - to match (3D-register) them across several stereo-pairs (visual odometry or sparse 3D mapping)
 - then typical interest-points & descriptors used (e.g., Harris & SIFT)

(a very short glimpse into)

Simultaneous Localization and Mapping (SLAM)

SLAM

chicken & egg problem

- build map (needs localization)
- while using map for localization

main idea

- exploit spatial relations to previously visited places (loop closing)
- to bound the cumulative error

SLAM

usually treated as two parts


- front-end
 - 2D/3D registration (and odometry)
 - plus related uncertainties as weights
 - sequentially plus loop-closures
- back-end
 - generic optimization of the weighted constraints

Relative Pose Estimation via Sensor Data Registration

Sensor Data Registration


problem:

- given two sensor data sets
- find parameters of transforms to spatially align them,
- i.e., 3-DoF (2D), 4-DoF (2.5D) or 6-DoF (3D)
- incl. uncertainty estimates for prob. localization & SLAM




Example: Scan Matching

based on LRF data



match to

vector map




grid (scan or map)

Example: image registration

e.g., using visual **features**

- **interest points**
 - to get a feasible subset of the whole image
- **descriptors** at the interest points
 - to solve the correspondence problem



(Natural) Image Features

visual features

- interest points
 - do not consider all pixels
 - but find "distinctive" locations
- descriptors
 - generate "unique" representation of the location
 - to allow matching across different images



very active field in computer vision

- for stereo, registration, ...
- speed / robustness trade-off very important

popular robotics combination:

- Harris corner (interest)
- SIFT or SURF (descriptor)

Loop Closing

Loop Closing

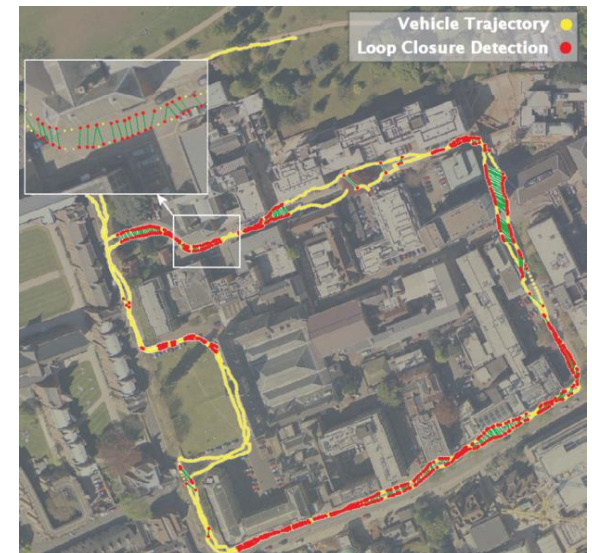

- proximity based
 - easiest possible strategy: use current localization estimate
 - to check whether there are previously visited places around
 - if so: try registrations
- place recognition
 - non-trivial, especially with respect to good strategies to get reasonable cost/benefit
 - typically (visual or 3D) feature collections in associative representation (hashes)
 - e.g., by using FABMAP in the context of metric SLAM

Appearance based Mapping: FAB-MAP

- learn a generative model of place appearance
- based on bag of words on SURF

very efficient

- sub-linear in the number of places



SLAM backend

SLAM backend:
Kalman filter

Kalman Filter for SLAM

- state vector robot motion in 2D
 - state vector is 3x1: [x,y,theta]
 - covariance matrix is 3x3
- mobile robot kinematics are not linear
 - use of EKF (or better UKF)
- SLAM: state vector is expanded
 - to include landmark positions
 - covariance matrix is also expanded

$$X = \begin{bmatrix} X_R^T & X_{L_1}^T & \dots & X_{L_n}^T \end{bmatrix}^T$$
$$\begin{bmatrix} P_{RR} & P_{RL_1} & \dots & P_{RL_N} \\ P_{L_1R} & P_{L_1L_1} & \dots & P_{L_1L_N} \\ \vdots & \vdots & \ddots & \vdots \\ P_{L_NR} & P_{L_NL_1} & \dots & P_{L_NL_N} \end{bmatrix}$$


Kalman Filter for SLAM

- challenges for Kalman Filter SLAM
 - n landmarks $\Rightarrow n^2$ variables in covariance
 - hence computationally challenging
(can be partially mitigated using sparse matrix algebra)
 - EKF tends to be unstable (hence option UKF)
- alternatives are hence popular
 - e.g., particle filter
 - e.g., graph based


SLAM backend: particle filter

Particle Filter

- represent distribution
 - by samples (particles)
 - population based
somewhat similar to
Evolutionary Algorithm
- (recap) example:
localization
 - (from Dieter Fox, UWash)
 - 24 sonar sensors
 - robot drawn at estimated
position,
 - which is not the correct
one in the beginning




Example Mapping



SLAM backend:
pose graph





Pose Graph

- nodes: sensor observations
 - e.g., 2D/3D scans, images
- edges: pose difference estimates with uncertainties
 - motion estimates via odometry and/or registration(s)
 - introduce constraints



Loop Closing

- use registration
- with previously visited places (pose-nodes)
- to improve map quality



Cost Function

- edge
 - motion estimates (pose differences)
 - with uncertainty (note: each edge assumed to be independent)
 - imagine spring/damper of certain length & stiffness
- Mahalanobis Distance
 - N-dim distance between x and y
 - including uncertainty in form of covariance C
 - respectively, its inverse aka information matrix Ω

$$x = (x_1, x_2, \dots, x_n)^T, \quad y = (y_1, y_2, \dots, y_n)^T$$

$$D_M(x, y) = \sqrt{(x - y)^T C^{-1} (x - y)}$$

Loop Closing & Relaxation

relaxation of the pose-graph

- find a configuration such that
 - the spring/damper network is
 - in an energetically optimal configuration
- i.e., minimize (least squares sense)
- sum of Mahalanobis Distances



That's all,
folks...