

Computer Networks - 2021

Mohammed El-Hajj

Jacobs University Bremen

Sept 3, 2021



JACOBS
UNIVERSITY



Computer Networks and Distributed Systems

- | | |
|---|------------------|
| • Introduction to Computer Science | 1st Semester |
| • Programming in C I | 1st Semester |
| • Algorithms and Data Structures | 2nd Semester |
| • Programming in C II | 2nd Semester |
| • Computer Architecture and Programming Languages | 3rd Semester |
| • Operating Systems | 3rd Semester |
| • Computer Networks | 4th Semester |
| • Distributed Algorithms | 6th Semester |
| • Project and Bachelor Thesis | 5th/6th Semester |

Credits

The following slides are the adaptation of:

- Lecture Slides of Prof [Jürgen Schönwälder](#) – Computer Networks 2019 – Jacobs University

Course Objectives

- Introduce fundamental data networking concepts
- Focus on widely deployed Internet protocols
- Prepare students for further studies in networking
- Combine theory with practical experiences
- Raise awareness of weaknesses of the Internet



Course Content

- 1.Introduction
- 2.Fundamental Networking Concepts
- 3.Local Area Networks (IEEE 802)
- 4.Internet Network Layer (IPv4, IPv6)
- 5.Internet Routing (RIP, OSPF, BGP)
- 6.Internet Transport Layer (UDP, TCP)
- 7.Firewalls and Network Address Translators
- 8.Domain Name System (DNS)
- 9.Abstract Syntax Notation 1 (ASN.1)
- 10.External Data Representation (XDR)
- 11.Augmented Backus Naur Form (ABNF)
- 12.Electronic Mail (SMTP, IMAP)
- 13.Document Access and Transfer (HTTP, FTP)

Reading Material

- A.S. Tanenbaum, "Computer Networks", 4th Edition, Prentice Hall, 2002
- W. Stallings, "Data and Computer Communications", 6th Edition, Prentice Hall, 2000
- C. Huitema, "Routing in the Internet", 2nd Edition, Prentice Hall, 1999
- D. Comer, "Internetworking with TCP/IP Volume 1: Principles Protocols, and Architecture", 4th Edition, Prentice Hall, 2000
- J.F. Kurose, K.W. Ross, "Computer Networking: A Top-Down Approach Featuring the Internet", 3rd Edition, Addison-Wesley 2004.

Grading

- Final examination (100%)
 - Covers the whole lecture
 - Closed book (and closed computers / networks)
- Quizzes (bonus) (5%)
 - Control your continued learning success
 - 3 quizzes with 10 pts each
 - 50 pts and above equals 30% of the overall grade
- Assignments(bonus) (5%)
 - Learning by solving assignments
 - Implement network protocols
 - Gain some practical experience in a lab session
 - 2 assignments with 10 pts each
 - 50 pts and above equals 30% of the overall grade

Part 1: Introduction

■ 1 Internet Concepts and Design Principles

■ Structure and Growth of the Internet

■ Internet Programming with Sockets

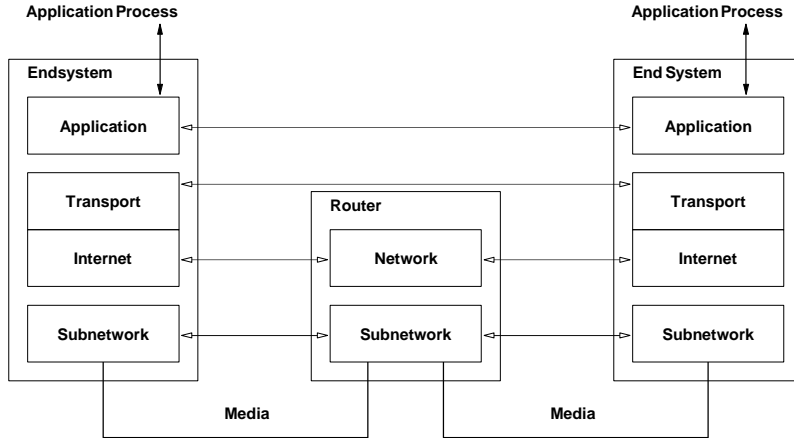
Section 1: Internet Concepts and Design Principles

1 Internet Concepts and Design Principles

Structure and Growth of the Internet

2 Internet Programming with Sockets

Internet Model



Internet Design Principles

- Connectivity is its own reward
- All functions which require knowledge of the state of end-to-end communication should be realized at the endpoints (end-to-end argument)
- There is no central instance which controls the Internet and which is able to turn it off
- Addresses should uniquely identify endpoints
- Intermediate systems should be stateless wherever possible
- Implementations should be liberal in what they accept and stringent in what they generate
- Keep it simple (when in doubt during design, choose the simplest solution)

Internet Design Principles

- Stateful services keep track of sessions or transactions and react differently to the same inputs based on that history.
- Stateless services rely on clients to maintain sessions and center around operations that manipulate resources, rather than the state.
- The internet moves from stateful to stateless architectures

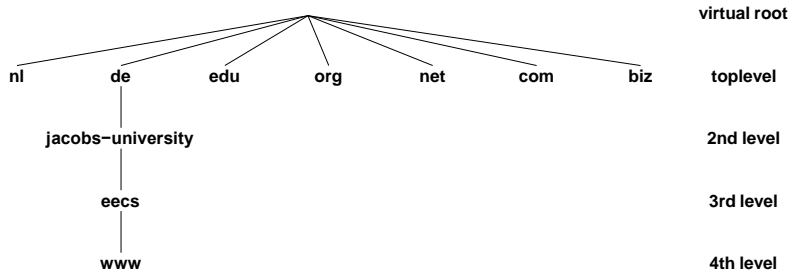
Internet Design Principles

- A communications network is a collection of systems (nodes) which transfer information between users attached to the networks.
- Two types of systems are defined within a communications network:
 - End Systems
 - Intermediate Systems

Internet Addresses

- Four byte IPv4 addresses are typically written as four decimal numbers separated by dots where every decimal number represents one byte (dotted quad notation). A typical example is the IPv4 address 192.0.2.1
- Sixteen byte IPv6 addresses are typically written as a sequence of hexadecimal numbers separated by colons (:) where every hexadecimal number represents two bytes
- Leading nulls in IPv6 addresses can be omitted and two consecutive colons can represent a sequence of nulls
- The IPv6 address 2001:00db8:0000:0000:0000:0000:0000:0001 can be written as 2001:db8::1
- See RFC 5952 for the recommended representation of IPv6 addresses

Internet Domain Names

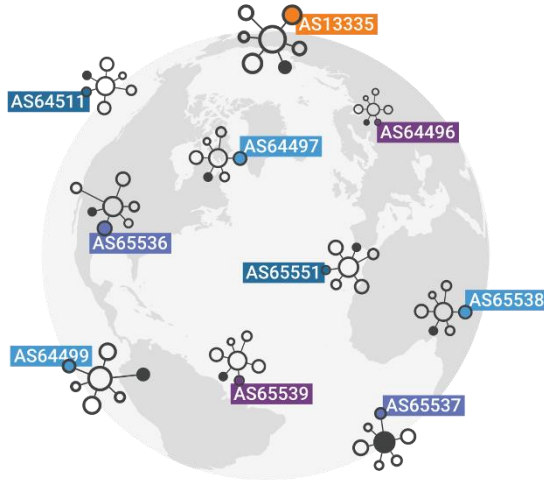


- The Domain Name System (DNS) provides a distributed hierarchical name space which supports the delegation of name assignments
- DNS name resolution translates DNS names into one or more Internet addresses

Autonomous Systems

- The global Internet consists of a set of inter-connected autonomous systems
- An *autonomous system* (AS) is a set of routers and networks under the same administration
- Autonomous systems are identified by 32-bit numbers, called AS numbers (ASNs) (originally the number space was limited to 16-bit but this has been increased to 32-bit)
- IP packets are forwarded between autonomous systems over paths that are established by an *Exterior Gateway Protocol* (EGP)
- Within an autonomous system, IP packets are forwarded over paths that are established by an *Interior Gateway Protocol* (IGP)

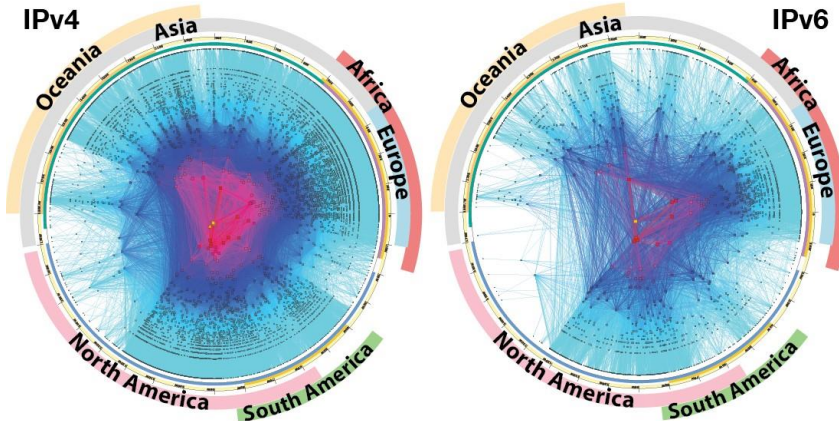
Autonomous Systems



Autonomous Systems

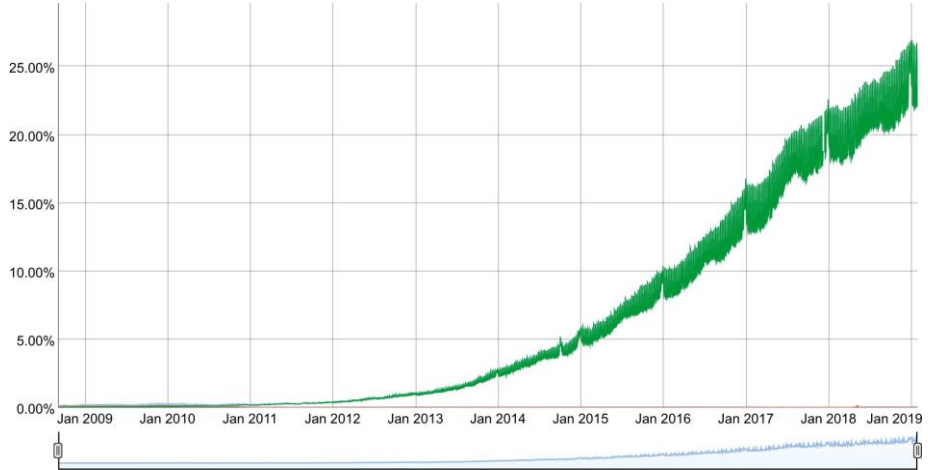
CAIDA's IPv4 vs IPv6 AS Core AS-level Internet Graph

Archipelago July 2015

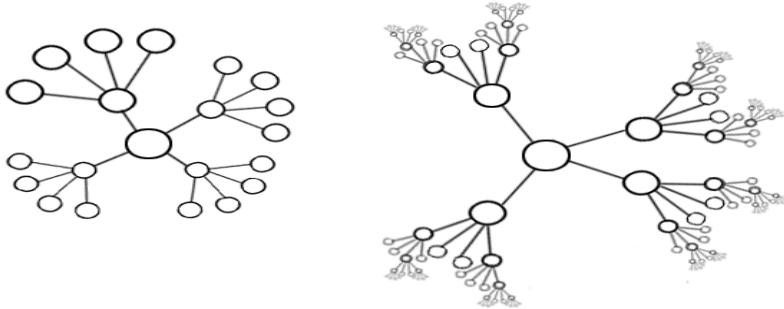


Copyright © 2015 UC Regents. All rights reserved.

Google IPv6 Adoption Statistics



Internet – A Scale-free Network?



- Scale-free: The probability $P(k)$ that a node in the network connects with k other nodes is roughly proportional to $k^{-\gamma}$ for some parameter γ .

Section 2: Structure and Growth of the Internet

1 Internet Concepts and Design Principles

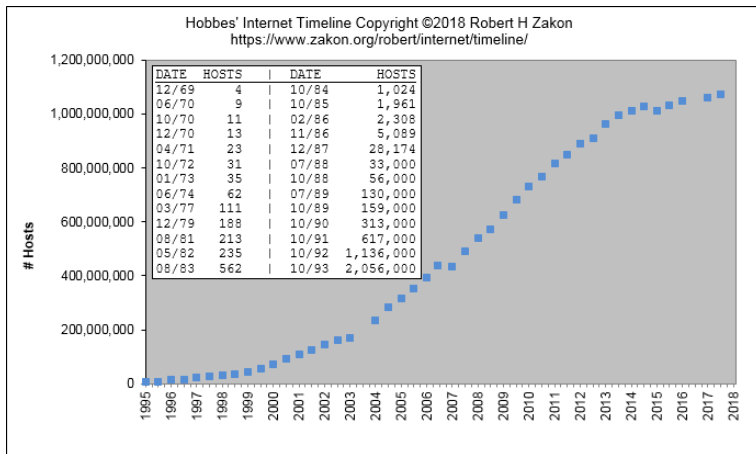
Structure and Growth of the Internet

Internet Programming with Sockets

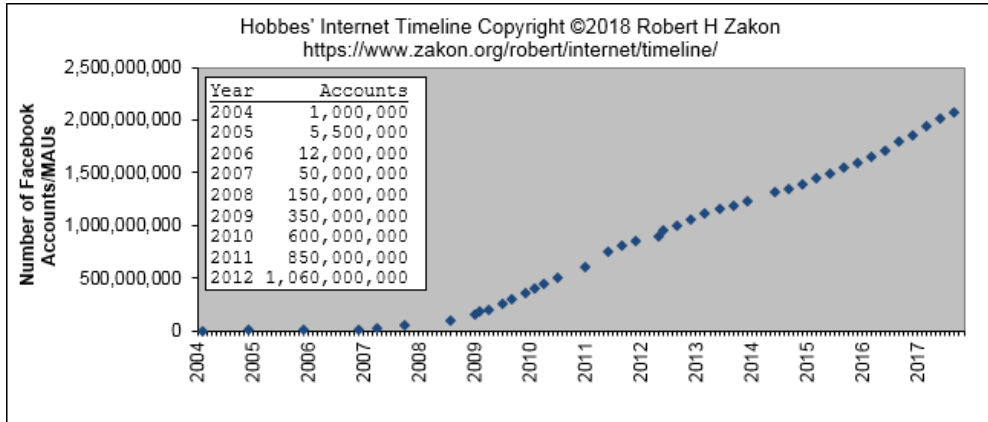
Evolution of Network Services

- 1970: private communication (email)
- 1980: discussion forums (usenet)
- 1985: software development and standardization (GNU)
- 1995: blogs, art, games, trading, searching (ebay, amazon)
- 1998: multimedia communication (rtp, sip, netflix)
- 2000: books and encyclopedia (wikipedia)
- 2005: social networks, video sharing (facebook, youtube)
- 2010: cloud computing, content delivery networks (akamai, amazon)
- 2015: voice-controlled personal assistants (amazon alexa, google home)
- 2020: Internet of Every Thing (IOE)

Growth of the Internet



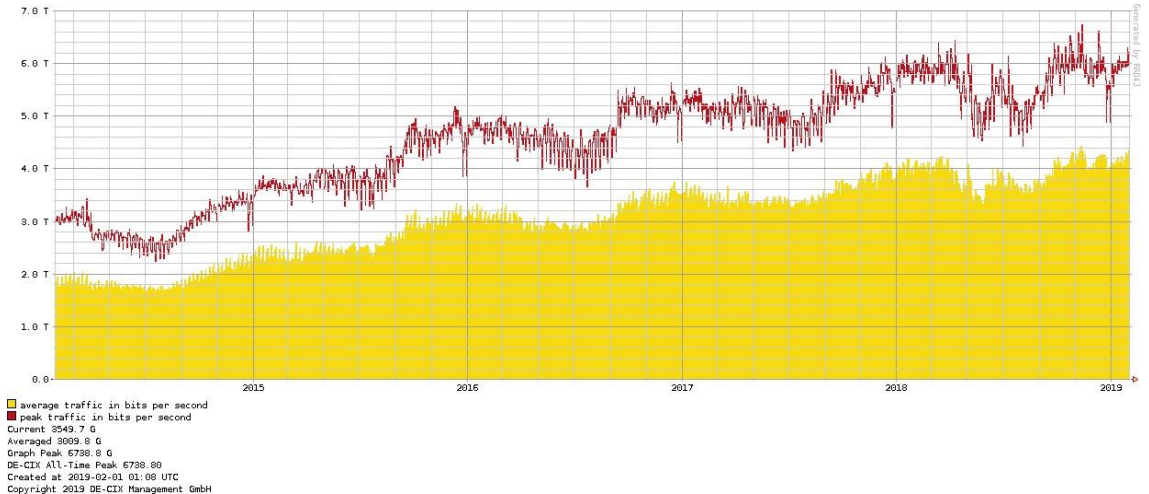
Growth of Facebook



Internet Exchange Points (2021)

- Internet Exchange Frankfurt/Germany (DE-CIX)
 - Connected networks: ≈ 800
 - Average throughput (1 year): $\approx 4.0\text{ Tbps}$
 - Maximum throughput (1 year): $\approx 6.0\text{ Tbps}$
 - Maximum transport capacity: $\approx 8\text{ Tbps}$
 - Total optical backbone capacity: $\approx 48\text{ Tbps}$
 - $3 \times 160\ 100$ Gigabit-Ethernetports
 - <http://www.decix.de/>
- Amsterdam Internet Exchange (AMS-IX)
 - <http://www.ams-ix.net/>
- London Internet Exchange (LINX)
 - <https://www.linx.net/>

DE-CIX Traffic (5 Years)



Networking Challenges

- Switching efficiency and energy efficiency
- Routing and fast convergence
- Security, trust, and key management
- Network measurements and automated network operations
- Ad-hoc networks and self-organizing networks
- Wireless sensor networks and the Internet of Things
- Delay and disruption tolerant networks
- High bandwidth and long delay networks
- Home networks, data center networks, access networks
- ...

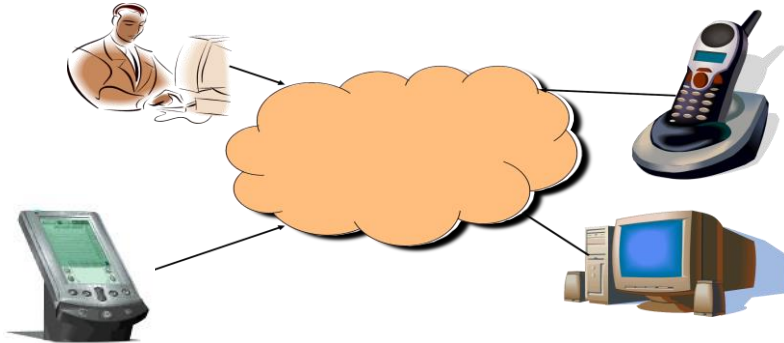
Section 3: Internet Programming with Sockets

- 1 Internet Concepts and Design Principles
- Structure and Growth of the Internet
- Internet Programming with Sockets

Internet Programming with Sockets

- Sockets are abstract communication endpoints with a rather small number of associated function calls
- The socket API consists of
 - address formats for various network protocol families
 - functions to create, name, connect, destroy sockets
 - functions to send and receive data
 - functions to convert human readable names to addresses and vice versa
 - functions to multiplex I/O on several sockets
- Sockets are the de-facto standard communication API provided by operating systems

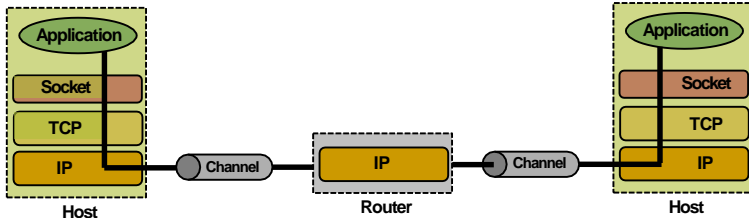
Internet Programming with Sockets



Also known as a "host"...

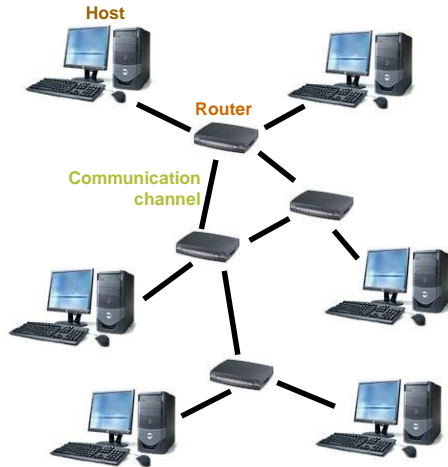
Berkley Sockets

- Universally known as **Sockets**
- It is an abstraction through which an application may send and receive data
- Provide **generic access** to interprocess communication services
 - ▣ e.g. IPX/SPX, Appletalk, TCP/IP
- Standard API for networking



Internet Programming with Sockets

- Computer Network
 - hosts, routers, communication channels
- **Hosts** run applications
- **Routers** forward information
- **Packets**: sequence of bytes
 - contain control information
 - e.g. destination host
- **Protocol** is an agreement
 - meaning of packets
 - structure and size of packetse.g. Hypertext Transfer Protocol (HTTP)



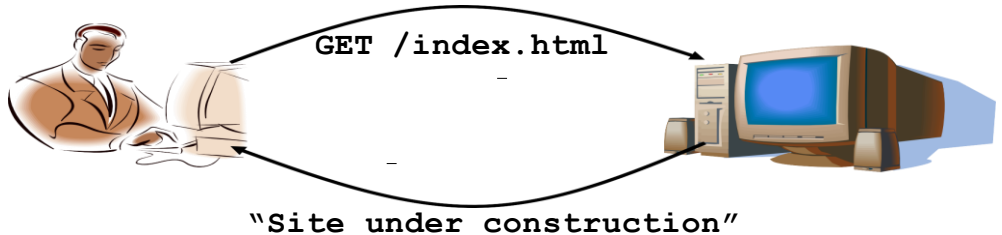
Client and Server

Client program

- Running on end host
- Requests service
- E.g., Web browser

Server program

- Running on end host
- Provides service
- E.g., Web server



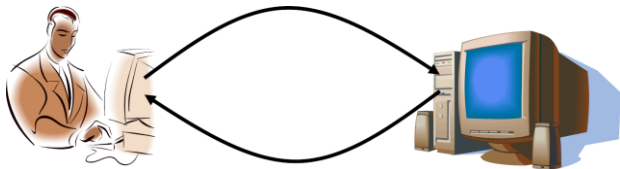
Client-Server Communication

Client

- Sometimes on
- Initiates a request to the server when interested
- E.g., web browser
- Needs to know the server's address

Server

- Always on
- Serve services to many clients
- E.g., <https://www.jacobs-university.de/>
- Not initiate contact with the clients
- Needs a fixed address



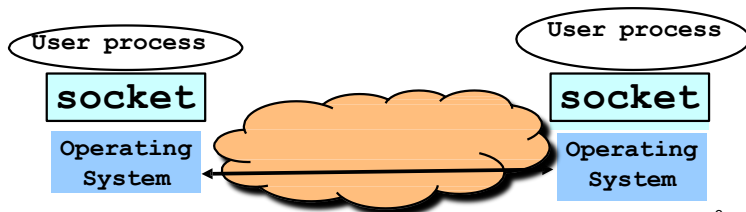
Socket: End Point of Communication

Processes send messages to one another

- Message traverse the underlying network

A Process sends and receives through a “socket”

- Analogy: the doorway of the house.
- Socket, as an API, supports the creation of network applications



UNIX Socket API

Socket interface

- A collection of system calls to write a networking program at user-level.
- Originally provided in Berkeley UNIX
- Later adopted by all popular operating systems

In UNIX, everything is like a file

- All input is like reading a file
- All output is like writing a file
- File is represented by an integer file descriptor
- Data written into socket on one host can be read out of socket on other host

System calls for sockets

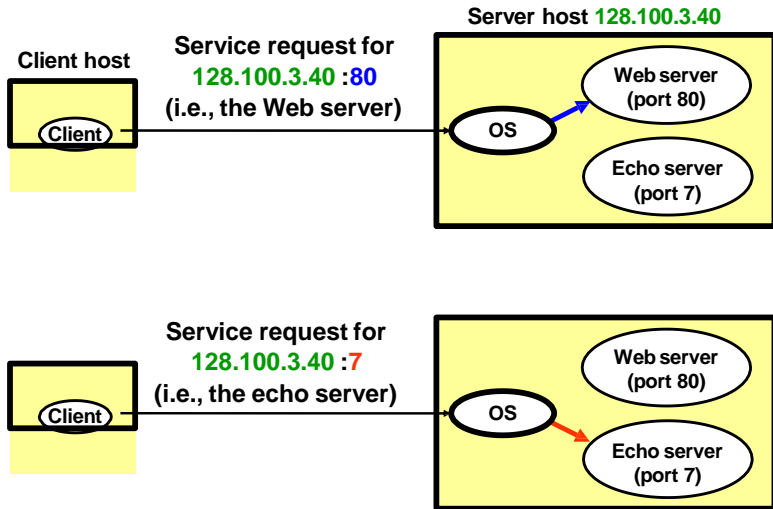
- Client: create, connect, write, read, close
- Server: create, bind, listen, accept, read, write, close

Typical Client Program

Prepare to communicate

- Create a socket
- Determine server address and port number
- Why do we need to have port number?

Using Ports to Identify Services



Socket Parameters

A socket connection has 5 general parameters:

- The protocol
 - Example: TCP and UDP.
- The local and remote address
 - Example: 128.100.3.40
- The local and remote port number
 - Some ports are reserved (e.g., 80 for HTTP)
 - Root access require to listen on port numbers below 1024

DNSlookup	UDP	53
FTP	TCP	21
HTTP	TCP	80
POP3	TCP	110
Telnet	TCP	23

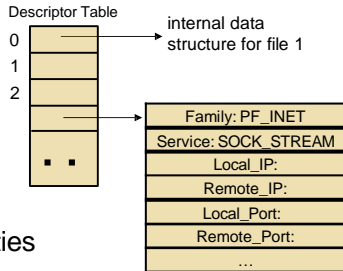
Sockets

■ Uniquely identified by

- an internet address
- an end-to-end protocol (e.g. TCP or UDP)
- a port number

■ Two types of (TCP/IP) sockets

- **Stream** sockets (e.g. uses TCP)
 - provide reliable byte-stream service
- **Datagram** sockets (e.g. uses UDP)
 - provide best-effort datagram service
 - messages up to 65.500 bytes
 - Socket extend the conventional UNIX I/O facilities
 - file descriptors for network communication
 - extended the read and write system calls



Typical Client Program

Prepare to communicate

- Create a socket
- Determine server address and port number
- Initiate the connection to the server

Exchange data with the server

- Write data to the socket
- Read data from the socket
- Do stuff with the data (e.g., render a Web page)

Close the socket

Important Functions for Client Program

- `socket()`
create the socket descriptor
- `connect()`
connect to the remote server
- `read()`, `write()`
communicate with the server
- `close()`
end communication by closing
socket descriptor

Creating a Socket

int socket(int domain, int type, int protocol)

- Returns a **descriptor** (or handle) for the socket
- **Domain**: protocol family
 - PF_INET for the Internet
- **Type**: semantics of the communication
 - SOCK_STREAM: Connection oriented
 - SOCK_DGRAM: Connectionless
- **Protocol**: specific protocol
 - UNSPEC: unspecified
 - (PF_INET and SOCK_STREAM already implies TCP)
- **E.g., TCP**: `sd = socket(PF_INET, SOCK_STREAM, 0);`
- **E.g., UDP**: `sd = socket(PF_INET, SOCK_DGRAM, 0);`

Connecting to the Server

- *int connect(int sockfd, struct sockaddr *server_address, socketlen_t addrlen)*
 - Arguments: socket descriptor, server address, and address size
 - Remote address and port are in struct sockaddr
 - Returns 0 on **SUCCESS**, and -1 if an error occurs

Sending and Receiving Data

Sending data

- *write(int sockfd, void *buf, size_t len)*
 - Arguments: socket descriptor, pointer to buffer of data, and length of the buffer
 - Returns the number of characters written, and -1 on error

Receiving data

- *read(int sockfd, void *buf, size_t len)*
 - Arguments: socket descriptor, pointer to buffer to place the data, size of the buffer
 - Returns the number of characters read (where 0 implies “end of file”), and -1 on error

Closing the socket

- *int close(int sockfd)*

Servers Differ From Clients

Passive open

- Prepare to accept connections
- ... but don't actually establish one
- ... until hearing from a client

Hearing from multiple clients

- Allow a backlog of waiting clients
- ... in case several try to start a connection at once

Create a socket for each client

- Upon accepting a new client
- ... create a *new* socket for the communication



Typical Server Program

Prepare to communicate

- Create a socket
- Associate local address and port with the socket

Wait to hear from a client (passive open)

- Indicate how many clients-in-waiting to permit
- Accept an incoming connection from a client

Exchange data with the client over new socket

- Receive data from the socket
- Send data to the socket
- Close the socket

Repeat with the next connection request

Important Functions for Server Program

- `socket()`
create the socket descriptor
- `bind()`
associate the local address
- `listen()`
wait for incoming connections from clients
- `accept()`
accept incoming connection
- `read(),write()`
communicate with client
- `close()`
close the socket descriptor

Socket Preparation for Server

Bind socket to the local address and port

- *int bind (int sockfd, struct sockaddr *my_addr, socklen_t addrlen)*
- Arguments: socket descriptor, server address, address length
- Returns 0 on success, and -1 if an error occurs

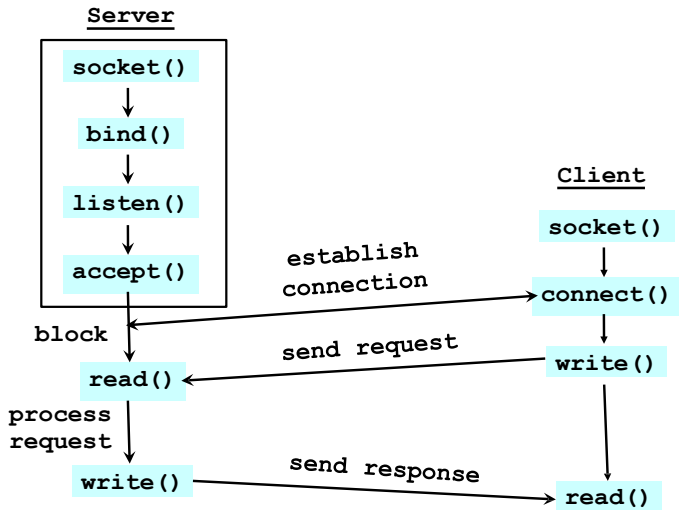
Define the number of pending connections

- *int listen(int sockfd, int backlog)*
- Arguments: socket descriptor and acceptable backlog
- Returns 0 on success, and -1 on error

Server Operation

- **accept()** returns a new socket descriptor as output
- New socket should be closed when done with communication
- Initial socket remains open, can still accept more connections

Putting it All Together



Supporting Function Calls

gethostname() get address for given host name (e.g. 128.100.3.40 for name “cs.toronto.edu”);

getservbyname() get port and protocol for a given service e.g. ftp, http (e.g. “http” is port 80, TCP)

getsockname() get local address and local port of a socket

getpeername() get remote address and remote port of a socket

IPv4 Socket Addresses

```
#include <sys/socket.h>
#include <netinet/in.h>

typedef ... sa_family_t;
typedef ... in_port_t;

struct in_addr { uint8_t
    s_addr[4];          /* IPv4 address */
};

struct sockaddr_in {
    uint8_t      sin_len;      /* address length (BSD) */
    sa_family_t sin_family;    /* address family */
    in_port_t    sin_port;     /* transport layer port */
    struct in_addr sin_addr;    /* IPv4 address */
};
```

IPv6 Socket Addresses

```
#include <sys/socket.h>
#include <netinet/in.h>

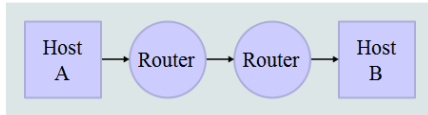
typedef ... sa_family_t;
typedef ... in_port_t;

struct in6_addr {
    uint8_t  s6_addr[16];          /* IPv6 address */
};

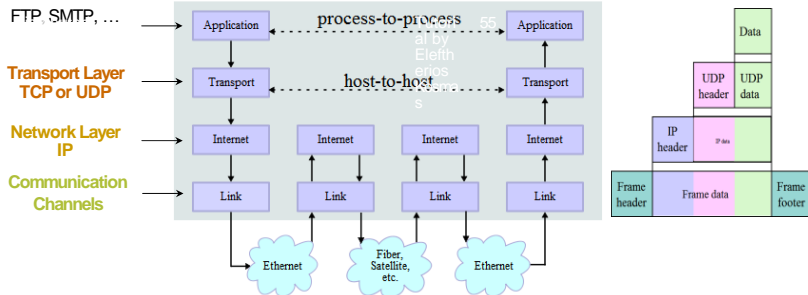
struct sockaddr_in6 {
    uint8_t    sin6_len;           /* address length (BSD) */
    sa_family_t sin6_family;       /* address family */
    in_port_t   sin6_port;         /* transport layer port */
    uint32_t    sin6_flowinfo;     /* flow information */
    struct in6_addr sin6_addr;     /* IPv6 address */
    uint32_t    sin6_scope_id;     /* scope identifier */
};
```

TCP/IP

Network Topology *

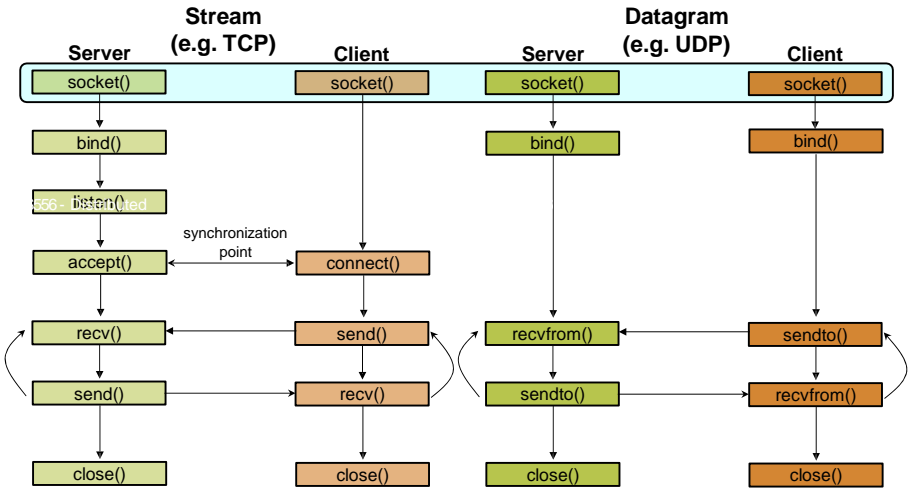


Data Flow



* image is taken from http://en.wikipedia.org/wiki/TCP/IP_model

Client - Server Communication - Unix



Socket API Summary

```
#include <sys/types.h>
#include <sys/socket.h>
#include <unistd.h>

#define          ...
SOCK_STREAM     ...
#define          ...
SOCK_DGRAM      ...
#define SCOK_RAW
#define SOCK_RDM
#define SOCK_SEQPACKET
...

#define AF_LOCAL ...
#define AF_INET  ...
#define AF_INET6 ...

#define PF_LOCAL ...
#define PF_INET  ...
#define PF_INET6 ...
```

Socket API Summary

```
int socket(int domain, int type, int protocol); int bind(int  
socket, struct sockaddr *addr,  
    socklen_t addrlen);
```

```
int connect(int socket, struct sockaddr *addr, socklen_t  
    addrlen);
```

```
int listen(int socket, int backlog);
```

```
int accept(int socket, struct sockaddr *addr, socklen_t  
    *addrlen);
```

```
ssize_t write(int socket, void *buf, size_t count); int send(int  
socket, void *msg, size_t len, int flags);
```

```
int sendto(int socket, void *msg, size_t len, int flags,  
    struct sockaddr *addr, socklen_t addrlen);
```

```
ssize_t read(int socket, void *buf, size_t count);
```

```
int recv(int socket, void *buf, size_t len, int flags); int recvfrom(int  
socket, void *buf, size_t len, int flags,  
    struct sockaddr *addr, socklen_t *addrlen);
```

Socket API Summary

```
int shutdown(int socket, int how); int  
close(int socket);
```

```
int getsockopt(int socket, int level, int optname, void *optval,  
               socklen_t *optlen);  
int setsockopt(int socket, int level, int optname,  
               void *optval, socklen_t optlen); int  
getsockname(int socket, struct sockaddr *addr,  
              socklen_t *addrlen);  
int getpeername(int socket, struct sockaddr *addr,  
                 socklen_t *addrlen);
```

- All API functions operate on abstract socket addresses
- Not all functions make equally sense for all socket types

Mapping Names to Addresses

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>

#define AI_PASSIVE ...
#define AI_CANONNAME ...
#define AI_NUMERICHOST ...

struct addrinfo { int
    int          ai_flags;
    int int      ai_family;
    size_t       ai_socktype;
               ai_protocol;
               ai_addrlen;
    struct sockaddr *ai_addr; char
               *ai_canonname;
    struct addrinfo *ai_next;
};
```

Mapping Names to Addresses

```
int getaddrinfo(const char *node,  
               const char *service,  
               const struct addrinfo *hints, struct  
               addrinfo **res);  
void freeaddrinfo(struct addrinfo *res);  
const char *gai_strerror(int errcode);
```

- Many books still document the old name and address mapping functions
 - gethostbyname()
 - gethostbyaddr()
 - getservbyname()
 - getservbyaddr()
- which are IPv4 specific and should not be used anymore

Mapping Addresses to Names

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>

#define NI_NOFQDN      ...
#define NI_NUMERICHOST ...
#define NI_NAMEREQD    ...
#define NI_NUMERICSERV ...
#define NI_NUMERICSCOPE ...
#define NI_DGRAM       ...

int getnameinfo(const struct sockaddr *sa,
                socklen_t salen,
                char *host, size_t hostlen, char
                *serv, size_t servlen, int flags);
const char *gai_strerror(int errcode);
```

Multiplexing

```
#include <sys/select.h>






typedef ... fd_set;

FD_ZERO(fd_set *set);
FD_SET(int fd, fd_set *set);
FD_CLR(int fd, fd_set *set);
FD_ISSET(int fd, fd_set *set);

int select(int n, fd_set *readfds, fd_set *writefds, fd_set
           *exceptfds, struct timeval *timeout);
int pselect(int n, fd_set *readfds, fd_set *writefds,
            fd_set *exceptfds, struct timespec *timeout, sigset_t
            sigmask);
```

- select() works with arbitrary file descriptors
- select() frequently used to implement the main loop of event-driven programs

References

- 
- B. Carpenter.
Architectural Principles of the Internet.
RFC 1958, IAB, June 1996.
- 
- B. Carpenter.
Internet Transparency.
RFC 2775, IBM, February 2000.
- 
- R. Gilligan, S. Thomson, J. Bound, J. McCann, and W. Stevens.
Basic Socket Interface Extensions for IPv6.
RFC 3493, Intrinsa Inc., Cisco, Hewlett-Packard, February 2003.
- 
- R. Atkinson and S. Floyd.
IAB Concerns and Recommendations Regarding Internet Research and Evolution.
RFC 3869, Internet Architecture Board, August 2004.
- 
- S. Kawamura and M. Kawashima.
A Recommendation for IPv6 Address Text Representation.
RFC 5952, NEC BIGLOBE, Ltd., NEC AccessTechnica, Ltd., August 2010.